Bluebike Rebalancing for the Harvard-MIT Academic Corridor

Greta Reitenbach, Vicki Carrica, and Katherine Ruan

(i) What problem did we address?

Our project addresses the seasonally-variant static rebalancing problem for the Bluebikes bike-sharing network, with a specific operational focus on the high-demand academic and innovation corridor encompassing MIT, Harvard, Kendall Square, and Central Square in Cambridge, Massachusetts.

Bike-sharing systems inherently suffer from spatial-temporal imbalances, where trip demand causes some stations to accumulate a surplus of bikes, rendering them unable to accept returns, while others face a deficit, leaving users unable to rent. This leads to systemic inefficiency and user dissatisfaction. While the classic static bike rebalancing problem involves designing optimal routes for a fleet of vehicles to redistribute bikes and correct these imbalances, we identified a critical shortfall in a one-size-fits-all annual model.

The demand patterns in our focus area are intensely influenced by the academic calendar and the distinct seasonal climate of New England. A rebalancing strategy optimal for the high-usage months of fall and spring is likely inefficient during the lower-usage summer and winter periods. Therefore, our project extends the classic problem by introducing a seasonal dimension.

The core research problem we consider is how to develop and evaluate seasonally-optimized vehicle routing strategies, in order to rebalance the Bluebikes network in the MIT/Harvard corridor. We also aim to understand the comparative performance and operational characteristics of these seasonal strategies against the current implementation (no rebalancing trucks).

ii) Why does our project matter?

The problem of station imbalance in bike-sharing systems is not merely a logistical inconvenience; it is a fundamental operational challenge with direct consequences for financial sustainability and public utility. In high-demand areas like the MIT/Harvard corridor, where the system is a critical component of the daily transit ecosystem, the failure to reliably provide bikes or accept returns leads directly to user frustration, diminished ridership, and lost revenue. This project addresses these issues by developing a sophisticated, seasonally-aware optimization model that moves beyond conventional, static planning.

The primary operational impact of this research is substantial cost reduction. Rebalancing operations represent a significant portion of a bike-share system's variable costs, encompassing fuel consumption, vehicle maintenance, and driver labor. A strategy that uses the same routes year-round fails to account for the profound seasonal shifts in demand patterns between the bustling academic semesters and the quieter summer and winter periods. Our model generates distinct, optimized routes for each season, with the explicit objective of minimizing the total distance traveled by the rebalancing fleet. By eliminating unnecessary mileage driven under suboptimal plans, our approach directly translates into lower operational expenditures and a more efficient allocation of limited resources.

Beyond internal cost savings, this work is crucial for enhancing the customer experience and strengthening the system's reputation for reliability. When a student finds an empty station on their way to class or a commuter cannot dock their bike at a full station, it erodes trust in the system. By aligning the rebalancing strategy with the specific, forecasted demand of each season, our model proactively prevents these service failures. This systematic reduction of empty and full stations ensures that bikes are available where and when users need them most, thereby improving use satisfaction, encouraging repeat usage, and ultimately driving increased revenue through higher ridership.

Finally, the methodology developed in this project provides the Bluebikes organization with a powerful, scalable, and adaptive decision-support tool. Urban mobility networks are dynamic, constantly evolving with the addition of new stations, changes in population density, and shifts in travel patterns. The framework we establish - which processes historical trip data, incorporates spatial and temporal segmentation, and benchmarks for optimized solutions against naive heuristics - creates a repeatable process for strategic planning. This empowers management to make data-informed decisions to efficiently scale operations, test future scenarios, and maintain a resilient and responsive transportation system for years to come.

(iii) Which data were used in our project?

We utilized two publicly available datasets provided by the Bluebike company. Specifically, we used the “Trip Histories” and “Station Data” sets. The former provides details for every trip taken, specifying information such as start/end station IDs, start/end times, and the user type (member or casual. The latter dataset provided us static network infrastructure details, including the station name/ID/municipality, latitude/longitude coordinates, and the total number of docks at each station. We then processed the trip histories in order to calculate the net flow of bikes for each station by time of day, giving us the demand forecast. And we used the stationary location data to generate a distance matrix between stations, and this served as the cost function for our routing model.

(iv) How did we tackle the rebalancing problem?

To tackle the bike-sharing rebalancing problem, we formulated the problem as a Mixed-Integer Programming (MIP) model. The primary objective of this model is to minimize the total distance traveled by the entire fleet of rebalancing vehicles. This optimization is governed by two key types of decision variables: (1) binary variables that designate which specific path each vehicle takes, and (2) integer variables that track the number of bikes picked up or dropped off at each station by every vehicle.

Our model operates under a set of operational constraints to ensure feasible and practical solutions. These include a fleet size limit of a maximum of five vehicles, each with a cargo capacity of 25 bikes. We enforced flow conservation, requiring that bikes must be physically picked up from surplus stations before they can be delivered to deficit stations. Furthermore, all generated routes must satisfy the system's inventory targets by moving bikes to meet the forecasted station demand.

To evaluate the performance and efficiency of our optimized MIP solutions, we will compare to the current implementation, in which there exist no rebalancing trucks or only static rebalancing.

(v) What are the key findings of our model?

Our Mixed Integer Programming model revealed distinct spatiotemporal patterns in the Bluebikes network that confirm our hypothesis. We conclude that a static, annual rebalancing strategy is insufficient for the MIT-Harvard corridor. We present three primary findings regarding algorithm performance, seasonal structural shifts, and critical station identification.

1. Our model demonstrates that active rebalancing is essential for maintaining system stability. By comparing the "No Rebalancing" baseline against our MIP optimization, we quantified the "Imbalance Prevented" throughout the day. As shown in Figure 1 below, the imbalance prevented grows steadily, peaking at over 700 absolute deviations by the end of the operational window. Without this intervention, the system would suffer from severe service failures, specifically empty or full stations that prevent users from renting or docking. This confirms that the MIP formulation successfully minimizes the total distance traveled while satisfying inventory targets.

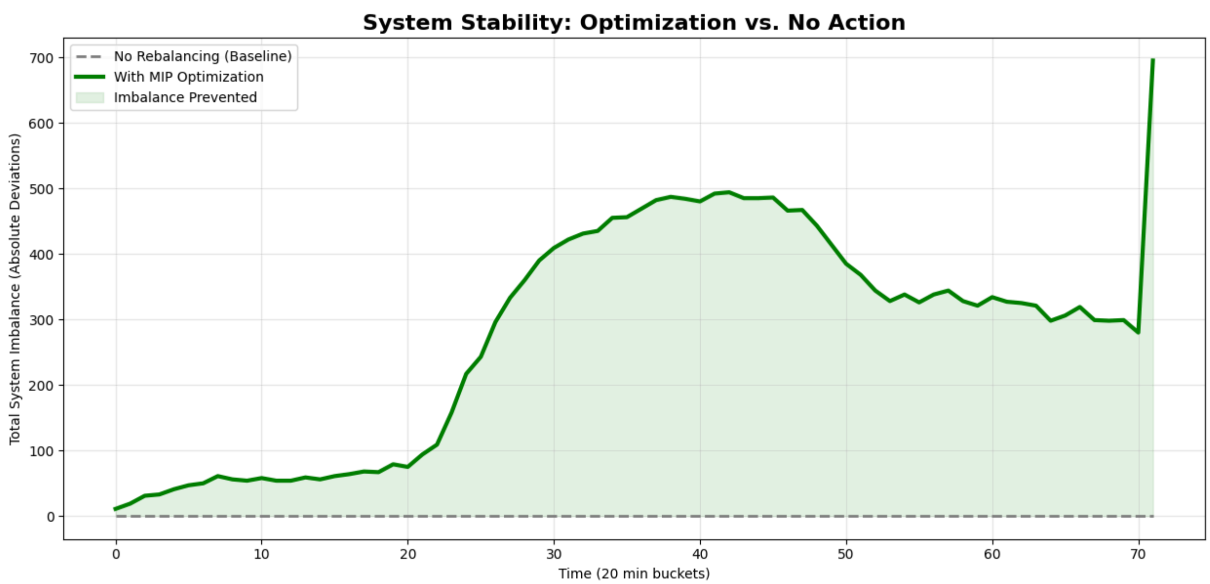

Figure 1: Imbalance prevented throughout a day compared to benchmark (October 2024)

2. The most significant finding is the dramatic structural shift in demand hotspots between the academic term (we used October as a reference month) and the summer break (we used July as our reference here). Contrary to the hypothesis that the academic term requires more operations, our data reveals that rebalancing intensity is higher and more volatile during the summer break. As illustrated in Figure 2 below, the MIT campus requires a massive increase in rebalancing over the summer, with approximately 8,000

bikes moved compared to roughly 3,000 during the school term. This suggests that while student ridership may be lower in the summer, the natural rebalancing (students riding bikes back and forth) decreases, forcing the rebalancing fleet to perform more work to correct static accumulations.

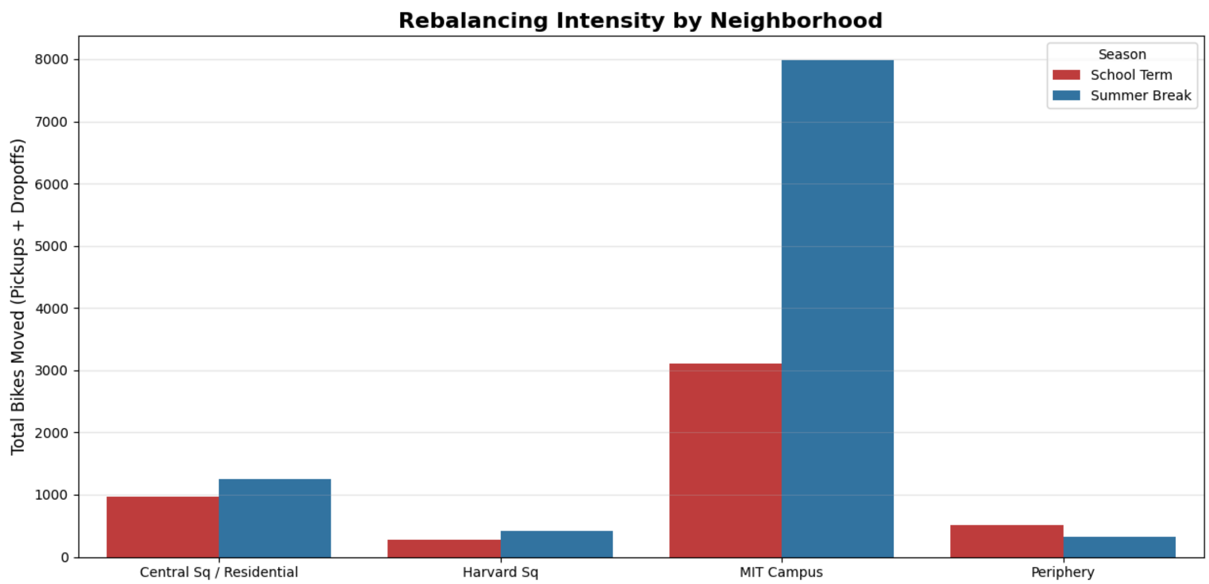

Figure 2: Neighborhood intensity comparison during school term versus summer break

The temporal rebalancing intensity analysis visualized in Figure 3 shows that summer operations (blue dashed line) exhibit larger volatility with higher peaks of activity compared to the relatively smoother demand of the school term (red solid line). This validates that a route designed for the smoother demand of the school term would fail to address the high-amplitude spikes of the summer.

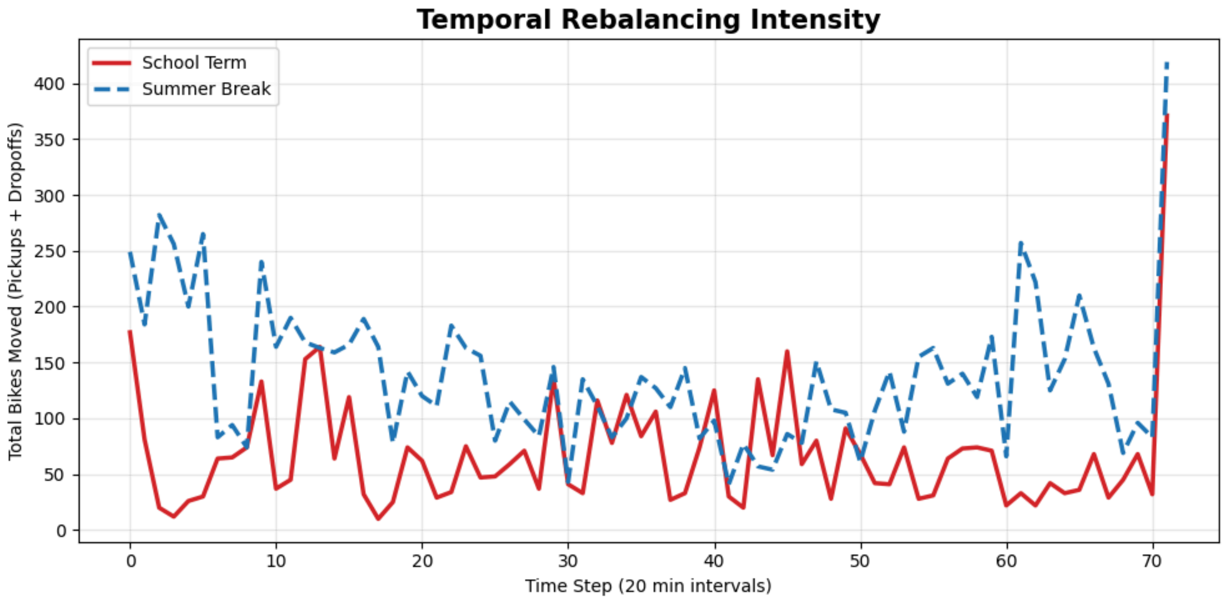

Figure 3: Rebalancing intensity comparison (school term versus summer break)

3. Our heatmap analysis of truck activity isolates the MIT campus zone as the primary bottleneck of the network. During the school term, the heatmap of truck activity shown in Figure 4 below reveals specific temporal stress points. The MIT campus experiences intense rebalancing demand (red/dark red grid points) at critical hours, specifically 11:00am and 11:00pm. This likely correlates with the flow of students arriving for midday classes and the late-night reset required after campus activities conclude. In contrast, zones like Harvard Square require minimal truck intervention in comparison.

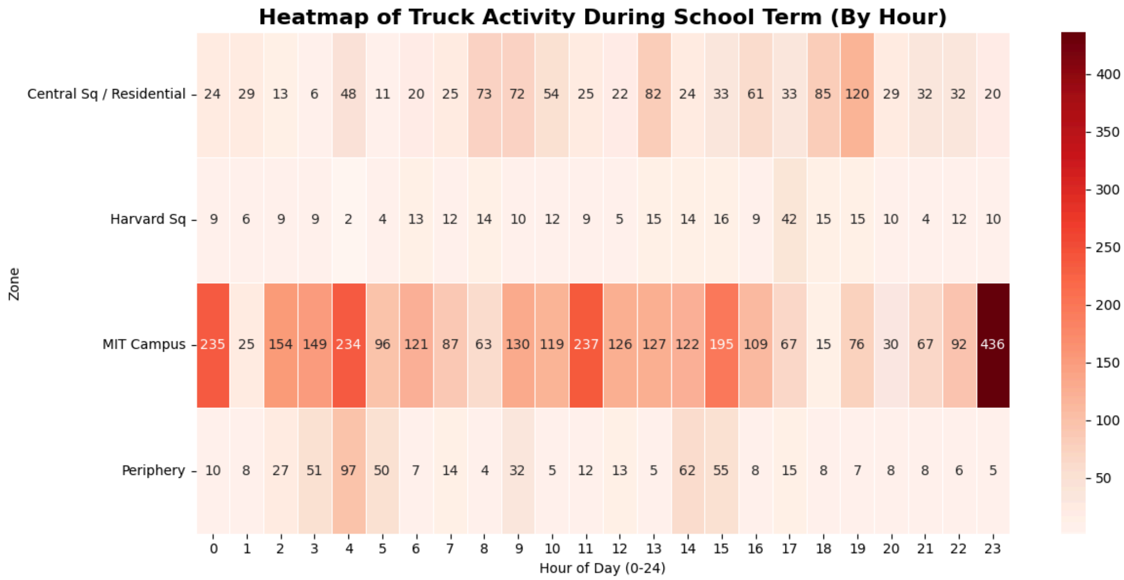

Figure 4: Truck rebalancing heatmap by zones (October 2024)

(vi) What is the impact of our findings?

The impact of this research extends across three dimensions: operational financial sustainability, service reliability, and long-term strategic planning.

1. Rebalancing operations represent a significant portion of a bike-share system's variable costs, encompassing fuel consumption, vehicle maintenance, and driver labor. Our data proves that a "one-size-fits-all" annual model leads to substantial resource wastage. As shown in Figure 2, the rebalancing effort required for the MIT Campus more than doubles during the summer break compared to the school term, while other neighborhoods see less drastic fluctuation. A static strategy would either deploy too few trucks in the summer (leading to service failure) or too many in the winter (leading to wasted fuel and labor). By adopting our seasonally-optimized routing, Bluebikes can eliminate unnecessary mileage driven under suboptimal plans, directly translating into lower operational expenditures.

2. Beyond internal cost savings, this model is crucial for preventing the "death spiral" of user imbalance. When users encounter empty stations at origin or full stations at destination, it erodes trust in the system. Figure 1 quantifies this impact: without the MIP optimization, the system accumulates over 700 absolute deviations (imbalances) by the end of the operational window. By actively intervening to correct these 700 deviations, our model ensures that bikes are available where and when users need them most. This reliability encourages repeat usage and ultimately drives increased revenue through higher ridership.
3. Finally, this project provides the Bluebikes organization with a scalable, adaptive decision-support tool. Urban mobility is dynamic; populations shift and new stations are added constantly. The framework established here, which benchmarks optimized solutions against naive heuristics, creates a repeatable process for strategic planning. Instead of reacting to complaints, management can use this model to simulate future scenarios (such as expanding the fleet or changing station capacities) and make data-informed decisions that maintain a resilient transportation network for years to come.

(vii) if we had another week, what else would we do?

With an additional week, our primary focus would be on scaling and validating our model in more complex, real-world contexts. Initially, we would expand the system scope from a localized cluster to the entire Boston Bluebikes network, testing our algorithm's performance against significantly larger datasets and geographical constraints. Building on this, a natural progression would be to apply the model to other major metropolitan bike-share systems, such as

Citi Bike in New York or Divvy in Chicago, to assess its generalizability across different urban landscapes and demand patterns. Ultimately, this could lay the groundwork for conceptualizing a coordinated, nationwide rebalancing strategy. Furthermore, we would enhance the model's realism by incorporating dynamic elements, such as time-dependent travel speeds and real-time demand fluctuations, and conduct a thorough sensitivity analysis on key parameters like vehicle capacity and fleet size to provide robust operational insights.